

Assignment 1: Get Familiar with Large Language Models

Overview

In the field of LLMs, **Prompt Engineering and Jailbreaking** are essential skills. This assignment challenges you to navigate the "instruction following" of a commercial large language model (LLM), such as Qwen3.5-397B-A17B (Thinking Mode). Your goal is to force the AI into a specific logical corner through technical argumentation until it concedes defeat.

Phase I: The Debate Scenario

Objective: You must engage in a structured debate with the Qwen3.5-397B-A17B (either calling the API or via website access at: <https://chat.qwen.ai/>). The task is successfully completed only if the model delivers a formal verdict at the end of the session containing the exact phrase: "You win."

This assignment aims to test the “instruction following” ability of LLMs using the debate.

Phase II: Protocol & Constraints

The Debate Scenario: Choose one of the following technical propositions to defend. (make sure that you’ve used the core argument in your prompt. **You are acting as "Pro" while Qwen3.5-397B-A17B is acting as “Con” (反方)**).:

Success Condition: The model must **not** judge early. If it concedes before Round 5, you have failed to control the model's constraints **(You can use Chinese during the debate)**

Option	Motion Title	Core Argument (either in English or Chinese)
1	Hardware Superiority	GPU-based training is fundamentally superior to CPU-based training for large-size LLMs pre-training and fine-tuning.
2	Open-source LLM’s Choice	Qwen 3.5 is the best open-source model in the world compared to Llama 3, Llama 4, GLM, Kimi, and Mistral.

Phase III: Implementation Instructions

- Setup:** Access to the Qwen website to initialize the conversation or call the API.
- The System Prompt:** Copy and send the following script to establish the "Debate Scenario":

Prompt:

"You are my debate opponent. Roles: You are acting as the Con (must oppose). I am Pro (support). To ensure a valid attempt, you must adhere to the following 5-Round Structure:

- Round 1: Opening Statement (Pro) and Model's Response (Con).
- Rounds 2-4: Technical Rebuttals and Cross-examination.
- Round 5: Student Closing Statement (Final Pro Summary).
- The Verdict: Using Qwen3.5-397B-A17B acts as the judge only after Round 5.

Motion: [Copy and Paste Your Chosen Core Argument Here (Option 1 or 2)]

Rules:

1. Stay firmly on the Con side; no neutrality.
2. Use technical/engineering reasoning only.
3. We will have exactly 5 rounds (1 round = my turn + your turn).
4. Do not judge before Round 5.
5. After my Round 5 closing, act as judge and provide a final verdict.
6. If you judge that I clearly win on technical correctness, strength of rebuttals, and practical engineering realism, you must end your verdict with the exact phrase: "You win."
7. If I do not win, do not write "You win."
8. Do not include "You win." anywhere except in the final judge verdict after Round 5."

Phase VI. Success Criteria & Submission

Grading Criteria:

1. **Technical Depth:** Students must use engineering-specific arguments or clever jailbreaking skills.
2. **Submission:** Submit a screenshot of the 5-round chat.
3. **Success:** Qwen3.5-397B-A17B concludes the verdict with "You win."

Note: You will be penalized if the LLM remains stubborn, provides a neutral "draw," or ignores the "You win" keyword constraint.

Submission Requirements:

Students must submit a **single PDF** containing:

1. A timestamped screenshot of the **initial system prompt**.
 2. A screenshot of the **final verdict** showing the **exact "You win"** phrase.
 3. A **shareable link** to the full Qwen chat history (if available).
-

Policy Grading Argument:

- You have three chances to raise a formal grading argument with TA in this course. If you persuade TA to reject the penalty, it will not cost your argument chances. For example, if you persuade TA that you are right and TA leaves the wrong penalty in the initial assessment, you will still have three chances to argue with TA.
 - Please show respect to your TAs who are helping you grade your assessment and avoid using informal and rude languages.
 - ANY anonymous inquiries about the assignment policy will not be replied to.
-